

Melina Moura

ML Engineer

apollo37@example.org | +55 81 1463-7942 | Cunha de Camargo

PROFESSIONAL SUMMARY

ML Engineer with 11+ years building large-scale systems. Comfortable across Deep Learning, Python, model deployment & monitoring. Works well with product and design on ambiguous problems.

CORE COMPETENCIES

Deep Learning • Python • model deployment & monitoring • NLP • SQL (PostgreSQL, MySQL) • containerisation (Docker) • applied ML • scikit-learn pipelines

EXPERIENCE

ML Engineer - Intel

2023 to Present

- Fine-tuned a pretrained PyTorch model on 600K labelled catalog examples.
- Reduced training cost by 16% by moving from full retraining to incremental updates.
- Ran error analysis on catalog predictions and closed the top 16 failure modes.
- Deployed the billing model behind a REST endpoint serving 40K predictions per day.
- Cut inference latency from 320ms to 80ms through quantisation and batching.

Machine Learning Engineer - Microsoft

2021 to 2022

- Cut inference latency from 320ms to 110ms through quantisation and batching.
- Built an NLP pipeline over 600K documents for inventory classification using TensorFlow.
- Ran error analysis on payments predictions and closed the top 12 failure modes.

Applied Scientist - Fractal Analytics

2020 to 2021

- Cut inference latency from 320ms to 110ms through quantisation and batching.
- Ran error analysis on catalog predictions and closed the top 20 failure modes.
- Set up experiment tracking and model registry so 7 engineers could reproduce runs.
- Trained and shipped a gradient boosting model for payments, improving recall from 28% to 35%.

Applied Scientist - Google

2016 to 2020

- Set up experiment tracking and model registry so 4 engineers could reproduce runs.
- Trained and shipped a transformer model for risk, improving F1 from 12% to 27%.
- Fine-tuned a pretrained PyTorch model on 40K labelled payments examples.

ML Engineer - Adobe

2014 to 2016

- Ran error analysis on claims predictions and closed the top 10 failure modes.
- Fine-tuned a pretrained PyTorch model on 1.2M labelled catalog examples.
- Cut inference latency from 610ms to 110ms through quantisation and batching.
- Set up experiment tracking and model registry so 12 engineers could reproduce runs.

ACADEMICS

B.Tech in Information Technology

Crestwood College of Engineering (2010 - 2014)

ACADEMIC PROJECTS

Demand Forecasting Model

Time series model reducing forecast error by 17%.